

# Adaptive Feature Selection in Random Forests

Christian Hilscher

University of Mannheim

Update on April 3, 2021

# How to build a tree 1/3

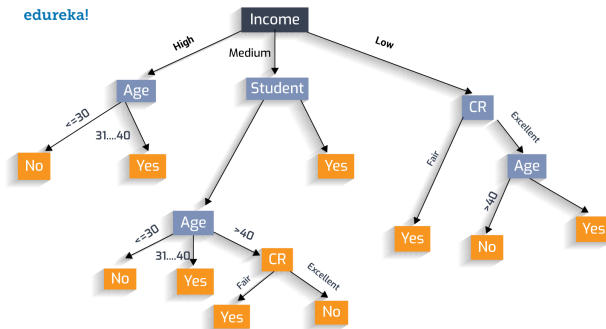


Figure: Example of a Decision Tree

Two questions at each node:

- ▶ Which variable to choose
- ▶ Which split-point to choose

## How to build a tree - choosing the split-point 2/3

- ▶ Pick one variable (= feature, dimension) and sort the X
- ▶ Each node has a certain variance:

$$\hat{\Delta}(t) = \frac{1}{N(t)} \sum_{\mathbf{x}_i \in t} (Y_i - \bar{Y}_i)^2$$

- ▶ Choose split-point such that the two daughter nodes are homogeneous
- ▶ This yields the following *decrease in impurity*:

$$\hat{\Delta}(s; \mathbf{t}) = \hat{\Delta}(\mathbf{t}) - [\hat{P}(\mathbf{t}_L)\hat{\Delta}(\mathbf{t}_L) + \hat{P}(\mathbf{t}_R)\hat{\Delta}(\mathbf{t}_R)]$$

- ▶ Maximising  $\hat{\Delta}(s; \mathbf{t})$  gives us split-point

## How to build a tree - choosing the variable 3/3

- ▶ Given a variable  $j$  we get  $\hat{\Delta}_j(s; \mathbf{t})$
- ▶ We then simply choose the variable with the highest  $j$ :

$$\arg \max_j \hat{\Delta}_j(s; \mathbf{t})$$

This continues until there are  $k \geq 1$  observation in each node

Estimator is then average in a node:

$$\hat{Y}_t = \frac{\sum_i Y_i \mathbb{1}_{(X_i \in t)}}{\sum_i \mathbb{1}_{(X_i \in t)}}$$

## Relevant concepts

Partial dependence function:

$$\bar{F}_j(x_j; t) = \mathbb{E}[Y | \mathbf{X} \in \mathbf{t}, X_j = j]$$

Mean-centered partial dependence function:

$$\begin{aligned}\bar{G}_j(x_j; t) &= \bar{F}_j(x_j; t) - \int \bar{F}_j(x_j; t) \mathbb{P}_{X_j | \mathbf{X} \in \mathbf{t}}(dx_j) \\ &= \mathbb{E}[Y | \mathbf{X} \in \mathbf{t}, X_j = j] - \mathbb{E}[Y | \mathbf{X} \in \mathbf{t}]\end{aligned}$$

Node balancedness:

$$\begin{aligned}\lambda_j(t) &= 4P_j(t_L)P_j(t_R) = 1 - |P_j(t_L) - P_j(t_R)|^2 \\ &= \frac{\Delta(j, s^*, t')}{|\bar{G}_j(x_j; t)|^2 + \Delta(j, s^*, t')}$$

# Splitting properties of the CART algorithm

- ▶  $\lambda_j(t) = \frac{\Delta(j, s^*, t')}{|\bar{G}_j(x_j; t)|^2 + \Delta(j, s^*, t')}$
- ▶ Then,  $\lim_{\Delta \rightarrow 0} \lambda = 0$  and  $\lim_{\Delta \rightarrow \infty} \lambda = 1$
- ▶ Noisy variables ( $\Delta \approx 0$ ) tend to be split on edges
- ▶  $\implies$  end-cut-preference (*ecp*)

Problems with splits on the edge:

- ▶ Very few observations in one of the two nodes
- ▶ Those nodes are not a good predictor

## Elevator Pitch

- ▶ We know that  $\arg \max_s \Delta$  leads to *ecp*
- ▶ Could instead maximise something else, for example

$$\Delta^\alpha = [4P(t_L)P(t_R)]^\alpha \Delta$$

- ▶ Since  $\Delta^\alpha \leq \Delta$ , this will lead to a lower decrease in impurity

## Elevator Pitch

- ▶ We know that  $\arg \max_s \Delta$  leads to *ecp*
- ▶ Could instead maximise something else, for example

$$\Delta^\alpha = [4P(t_L)P(t_R)]^\alpha \Delta$$

- ▶ Since  $\Delta^\alpha \leq \Delta$ , this will lead to a lower decrease in impurity

My idea:

- ▶ Our problem is essentially

$$\arg \max_j \arg \max_s \Delta_j^\alpha(s)$$

- ▶ Noisy variables are split at ends of feature space
- ▶ They get a low weight
- ▶ They will be chosen less often as optimal split dimension



# Questions

Two approaches:

$$\text{a.) } \arg \max_j \arg \max_s [4P(t_L)P(t_R)]^\alpha \Delta_j(s)$$

$$\text{b.) } \arg \max_j [4P(t_L)P(t_R)]^\alpha \arg \max_s \Delta_j(s)$$

a.) weights incorporated into choosing the split-point *within* a variable

b.) original CART with weights slapped on for penalizing *ecp*

## Roadmap - next 14 days

- ▶ Show that noisy variables are less often chosen with b.)  
(should be fairly ok, no?)
- ▶ Think about whether one could show the same with a.)  
(is this even true?)
- ▶ upgrade DTRegressor to RFR
- ▶ Play around and get maybe a feeling on how stuff changes empirically